

Computer Networks

Chapter 7: Application Layer (2/3)

(Version February 27, 2023)

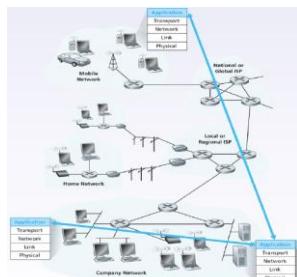
Xiao Mingzhong

CIST, Beijing Normal University

应用层

- Domain Name System(DNS)
- Electronic Mail*
- FTP*
- World Wide Web(WWW)
- Multimedia*

- Peer-to-Peer应用
 - Napster,Maze
 - Gnutella,BitTorrent
 - DHT(Chord,Kademlia)



Review

- 从E2E的传输抽象
- 到
- 使用TCP/UDP的Socket 程序设计
 - Tcp/udp是个什么样的服务
 - 如何使用它们? Socket程序设计.

- Now, 是琢磨应用及其协议的时候了.



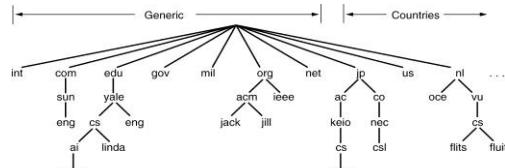
DNS 重要性

- Motivation(动机)
 - 通过网络地址直接访问资源，地址记忆难
 - 机器名与机器地址分离的需要
- Host.txt解决策略
 - 某机器记录所有主机的IP地址及其对应名字
 - 网络主机周期下载此文件到本地，以实现<名,址>转换
- 存在扩展性问题
 - DNS解决方案



域名系统(DNS)

- Basic idea: 每台主机都有一个全球唯一的名字跟它的IP地址关联, 由DNS系统来提供按名字查找IP地址的服务. 即: 提交一个主机名, 该系统返回该主机的IP地址. 命名?(层次命名法, 以net.bnu.edu.cn为例)

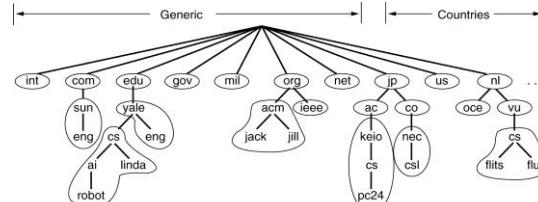


- 从叶子节点向上到达树根, 一路所遇节点名串联起来就是该主机的名字, 称为**域名**, 用“.”分隔. 一个域在域名空间树中就是一棵**子树**.
- 每个域应该部署有相应的**名字服务器(name server)**, 来负责存储和管理其下游名字跟IP地址的对应关系. 这种对应关系通常以文件形式存储, 专业术语资源记录(**resource records**)



DNS名字服务器

- 层次表达的名字空间被划分成一系列非重叠的区域(zones), 每个区域由一个或多个**name servers**管理, 该NS由其上一个NS管理:



- Note: 之所以多个服务器, 备份作用. “从”服务器通常需要跟“主”服务器保持数据一致, “主”服务器上的信息由域名管理者维护.
- 名字解析器(resolver) 负责向名字服务发起DNS请求. A resolver 通常以若干库例程(library routines)的方式被“链接”(linked)进应用程序中. 如: gethostbyname(...);



资源记录(Resource records)

→

Type	Associated entity	Description
SOA	Zone	Holds information on the represented zone
A	Host	Contains an IP address of the host this node represents
MX	Domain	Refers to a mail server to handle mail addressed to this node
SRV	Domain	Refers to a server handling a specific service
NS	Zone	Refers to a name server that implements the represented zone
CNAME	Node	Symbolic link with the primary name of the represented node
PTR	Host	Contains the canonical name of a host
HINFO	Host	Holds information on the host this node represents
TXT	Any kind	Contains any entity-specific information considered useful



Resource Record (Example)

```
; Authoritative data for cs.vu.nl
cs.vu.nl.      86400 IN SOA star boss (952771.7200.7200.2419200.86400)
cs.vu.nl.      86400 IN TXT "Divisie Wiskunde en Informatica."
cs.vu.nl.      86400 IN TXT "Vrije Universiteit Amsterdam."
cs.vu.nl.      86400 IN MX 1 zephyr.cs.vu.nl.
cs.vu.nl.      86400 IN MX 2 top.cs.vu.nl.

flits.cs.vu.nl. 86400 IN HINFO Sun Unix
flits.cs.vu.nl. 86400 IN A 130.37.16.112
flits.cs.vu.nl. 86400 IN A 192.31.231.165
flits.cs.vu.nl. 86400 IN MX 1 flits.cs.vu.nl.
flits.cs.vu.nl. 86400 IN MX 2 zephyr.cs.vu.nl.
flits.cs.vu.nl. 86400 IN MX 3 top.cs.vu.nl.
www.cs.vu.nl.   86400 IN CNAME star.cs.vu.nl
ftp.cs.vu.nl.   86400 IN CNAME zephyr.cs.vu.nl

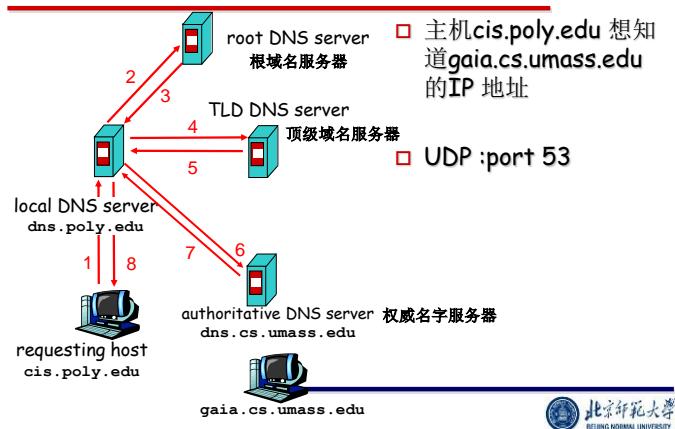
rowboat        IN A          130.37.56.201
                IN MX         1 rowboat
                IN MX         2 zephyr
                IN HINFO     Sun Unix

little-sister   IN A          130.37.62.23
                IN HINFO     Mac MacOS

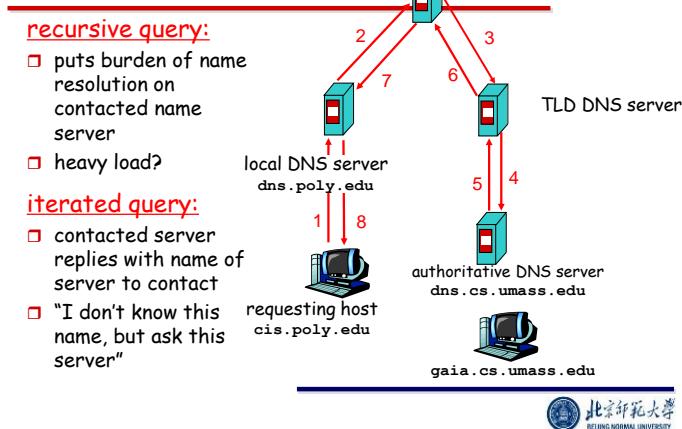
laserjet       IN A          192.31.231.216
                IN HINFO     "HP Laserjet IIIS" Proprietary
```



DNS 的迭代名字解析过程



递归解析过程

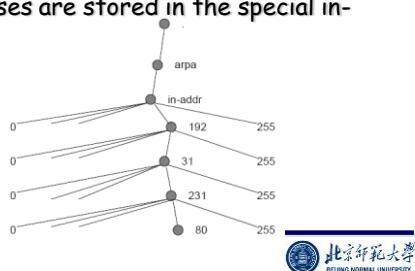


DNS Address Resolution(地址解析)

- Guess what: you can also find the name of a host when given its address:

```
% host 192.31.231.80
% 80.231.31.192.in-addr.arpa domain name pointer veldersschuit.cs.vu.nl.
```

- Solution: IP addresses are stored in the special in-addr.arpa domain:



Round robin DNS(轮回 DNS)

- 为地理上分布的多个服务器进行负载平衡用。
- 原理是针对DNS查询返回的不是一个IP地址,而是几个。

- e.g., executing host www.google.com returns

- www.google.com is an alias for www.l.google.com.
- www.l.google.com has address 66.249.91.99
- www.l.google.com has address 66.249.91.104
- www.l.google.com has address 66.249.91.103
- www.l.google.com has address 66.249.91.147

- Each DNS reply has associated a TTL(生存期), telling how long it can be cached.

- 某些情况下,返回的IP地址跟请求者的地理位置有关。(e.g., Akamai) 出于就近获得服务的目的,前面是负载平衡目的。
- Lab tools: nslookup & dig

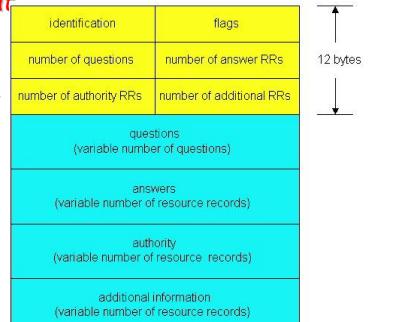


DNS protocol messages

DNS protocol : query and reply messages, both with same message format

msg header

- **identification:** 16 bit # for query, reply to query uses same #
- **flags:**
 - ❖ query or reply
 - ❖ recursion desired
 - ❖ recursion available
 - ❖ reply is authoritative



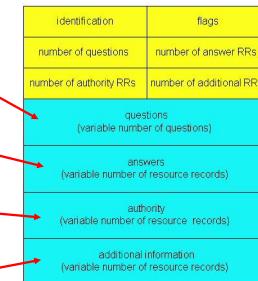
DNS protocol, messages

Name, type fields
for a query

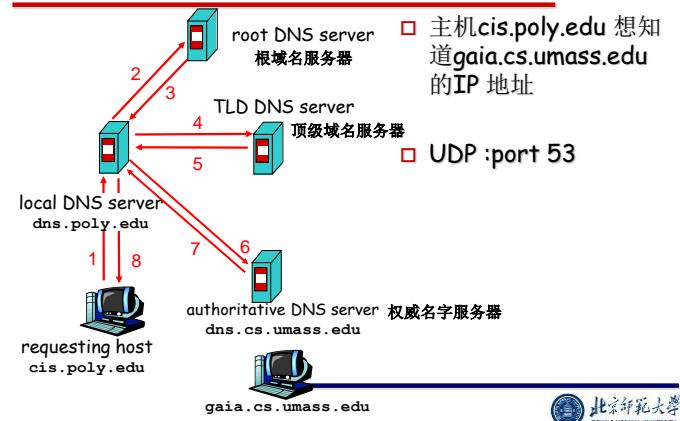
RRs in response
to query

records for
authoritative servers

additional "helpful"
info that may be used

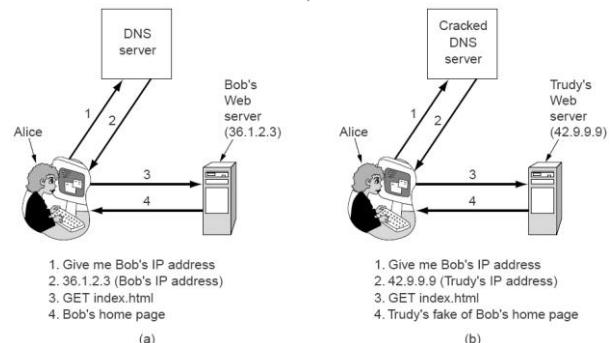


DNS 的迭代名字解析过程(回顾)



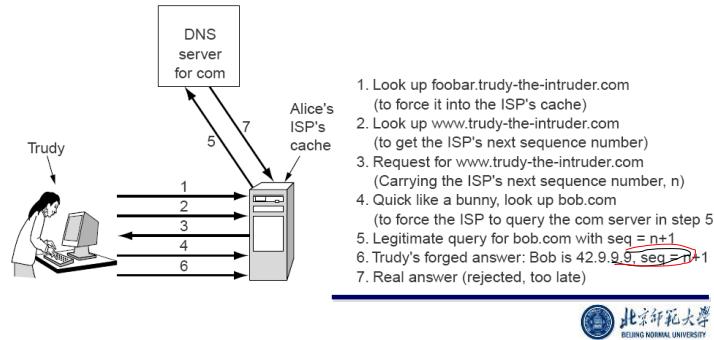
DNS 攻击例

□ Essence: 寻找DNS协议的问题点,让用户解析器获得错误的IP地址.



Steps:

- Observation: DNS 根据 UDP 包中的 "identification" 来关联"请求和应答", Trudy 能发起 DNS 骗骗:
 - 假定: DNS cache 初始为空(or cached entry 已过期)



Web App. and HTTP Pro.

First some jargon

- 一个 Web page 由若干 objects 构成
- Object 可能是 HTML file, JPEG image, Java applet, audio file, Link, ...
- Web page 的 objects 用 HTML 语言来“标记”
- Each object 用一个 URL 标识, 它是可寻址的
- Example URL:

`www.someschool.edu/someDept/pic.gif`

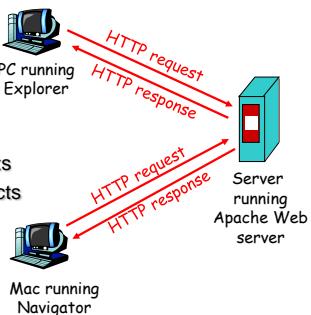
host name path name



HTTP overview

HTTP: hypertext transfer protocol 超文本传输协议

- Web's 应用层协议
- client/server model
 - **client**: browser that requests, receives, "displays" Web objects
 - **server**: Web server sends objects in response to requests
- HTTP 1.0: RFC 1945
- HTTP 1.1: RFC 2615



HTTP overview (continued)

Uses TCP:

- Client 发起对服务器 80 端口的 TCP connection 请求
- Server 接受 TCP connection 请求
- HTTP messages (应用层协议消息) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is "stateless" 无状态的

- server 不维护 client 的请求史

Protocols that maintain "state" are complex!

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled



HTTP connections

Nonpersistent HTTP 非持续的http连接

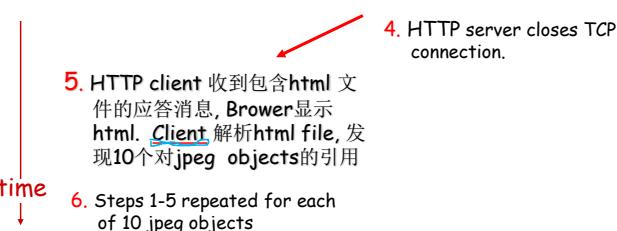
- At most one object is sent over a TCP connection. Web 页一个对象一个TCP连接。
- HTTP/1.0 uses nonpersistent HTTP

Persistent HTTP 持续的http连接

- Multiple objects can be sent over single TCP connection between client and server. 一个TCP连接上可发多个对象。
- HTTP/1.1 uses persistent connections in default mode



Nonpersistent HTTP (cont.)



Nonpersistent HTTP

Suppose user enters URL www.someSchool.edu/someDepartment/home.index

- (contains text, references to 10 jpeg images)
- 1a. HTTP client 发起TCP连接 to HTTP server (process) at www.someSchool.edu on port 80
 - 1b. HTTP server at host www.someSchool.edu waiting for TCP connection at port 80. "accepts" 连接, 并通知client
 - 2. HTTP client sends HTTP request message (containing URL) into TCP connection socket. Message indicates that client wants object someDepartment/home.index
 - 3. HTTP server receives request message, forms response message containing requested object, and sends message into its socket

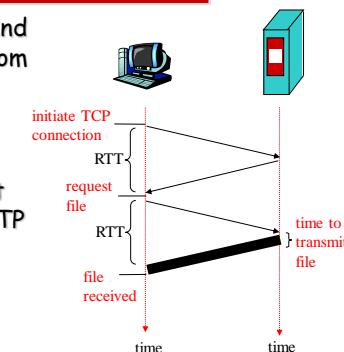


Response time modeling 响应时间模型

Definition of RRT: time to send a small packet to travel from client to server and back.

Response time:

- 一个RTT用于发起TCP连接
 - one RTT for HTTP request and first few bytes of HTTP response to return
 - file transmission time
- total = 2RTT+transmit time



持续HTTP

非持续HTTP分析:

- 要求2 RTTs per object
- 每次TCP连接都会引发OS开销
- browser often open parallel TCP connections to fetch referenced objects

持续HTTP

- 服务在响应请求后,不立即关闭连接
- subsequent HTTP messages between same client/server sent over open connection

无管道型持续http:

- 仅当前面的请求被响应, client才会发起新的请求.
- one RTT for each referenced object

管道型持续http:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects



HTTP request message

- two types of HTTP messages: **request, response**

HTTP request message:

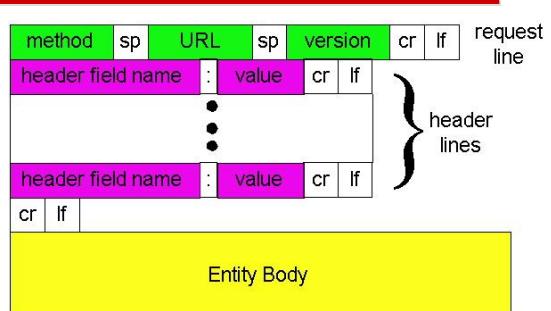
- ASCII (human-readable format)

request line
(GET, POST, HEAD commands)
header lines
Carriage return, line feed indicates end of message

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```



HTTP request message: general format



Uploading form input

Post method:

- Web page often includes form input
- Input is uploaded to server in entity body

URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

www.somesite.com/animalsearch?monkeys&banana



Method types

HTTP/1.0

- GET
- POST
- HEAD
 - is similar to GET, but asks server to leave requested object out of response

HTTP/1.1

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field



HTTP response message

status line
(protocol
status code
status phrase)
HTTP/1.1 200 OK
Connection close
header lines
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998
Content-Length: 6821
Content-Type: text/html

data, e.g.,
requested
HTML file
data data data data data ...



HTTP response status codes(应答状态码)

In first line in server->client response message.

A few sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

400 Bad Request

- request message not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported



Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

telnet cis.poly.edu 80

Opens TCP connection to port 80
(default HTTP server port) at cis.poly.edu.
Anything typed in sent
to port 80 at cis.poly.edu

2. Type in a GET HTTP request:

GET /~ross/ HTTP/1.1
Host: cis.poly.edu

By typing this in (hit carriage
return twice), you send
this minimal (but complete)
GET request to HTTP server

3. Look at response message sent by HTTP server!



User-server state: cookies

Many major Web sites
use cookies

Four components:

- 1) cookie header line of HTTP **response** message
 - 2) cookie header line in HTTP **request** message
 - 3) cookie **file** kept on user's host, managed by user's browser
 - 4) back-end **database** at Web site

■ One Note & Specifying commerce site for first time

■ When initial HTTP requests arrives at site, site creates a unique ID and creates an entry in backend database for ID

Example:

- Susan access Internet always from same PC
 - She visits a specific e-commerce site for first time
 - When initial HTTP requests arrives at site, site creates a unique ID and creates an entry in backend database for ID



Cookies (continued)

What cookies can bring:

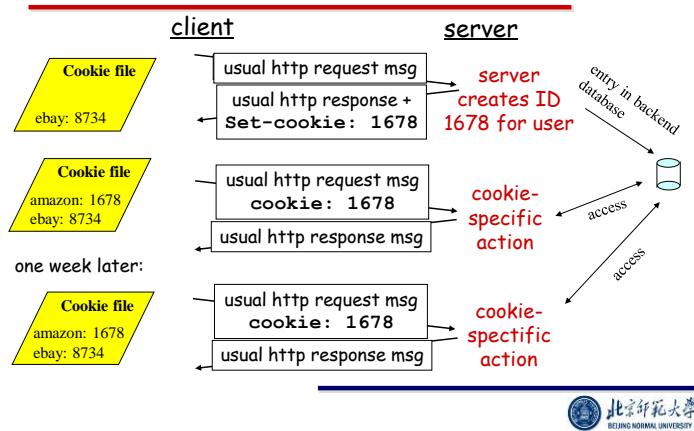
- authorization
 - shopping carts
 - recommendations
 - user session state
(Web e-mail)

Cookies and privacy:

- ❑ cookies permit sites to learn a lot about you
 - ❑ you may supply name and e-mail to sites
 - ❑ search engines use redirection & cookies to learn yet more
 - ❑ advertising companies obtain info across sites



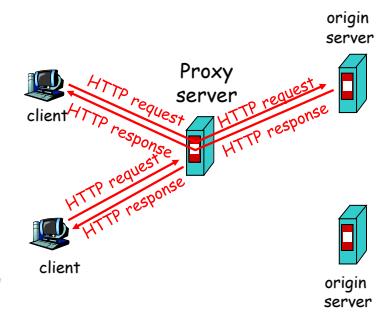
Cookies: keeping “state” (cont.)



Web caches (web proxy server)

Goal: satisfy client request without involving origin server

- ❑ user sets browser:
Web accesses via
cache
 - ❑ browser sends all
HTTP requests to
cache
 - object in cache: cache
returns object
 - else cache requests
object from origin server,
then returns object to
client



More about Web caching

- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)

Why Web caching?

- Reduce response time for client request.
- Reduce traffic on an institution's access link.
- Internet dense with caches enables "poor" content providers to effectively deliver content (but so does P2P file sharing)



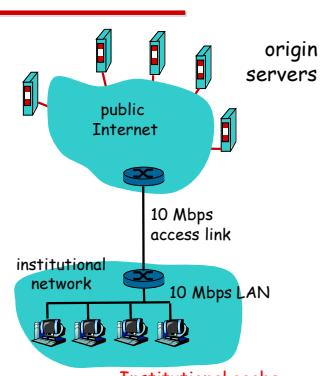
Caching example (cont)

Possible solution

- increase bandwidth of access link to, say, 10 Mbps

Consequences

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay
 $= 2 \text{ sec} + \text{msecs} + \text{msecs}$
- often a costly upgrade



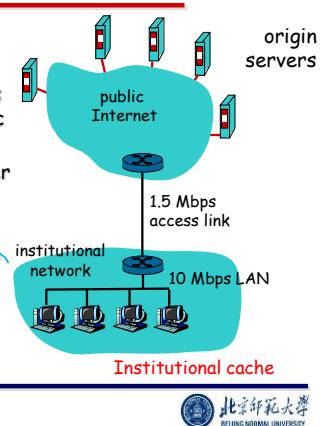
Caching example

Assumptions

- average object size = 100,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

Consequences

- utilization on LAN = 15% *(15 * 100k / 10M)*
- utilization on access link = 100% *(15 * 100k / 1.5M)*
- total delay = Internet delay + access delay + LAN delay
 $= 2 \text{ sec} + \text{minutes} + \text{milliseconds}$



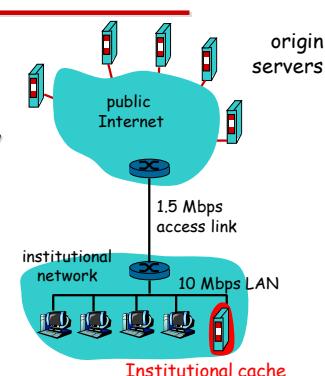
Caching example (cont)

Install cache

- suppose hit rate is .4

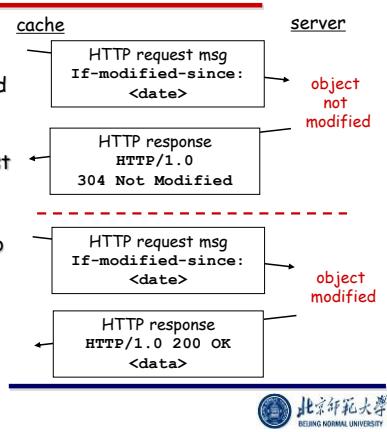
Consequence

- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + **access delay** + LAN delay
 $= .6 * (2.01) \text{ secs} + .4 \text{ milliseconds} < 1.4 \text{ secs}$



Conditional GET

- Goal: don't send object if cache has up-to-date cached version
- cache: specify date of cached copy in HTTP request
If-modified-since: <date>
- server: response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



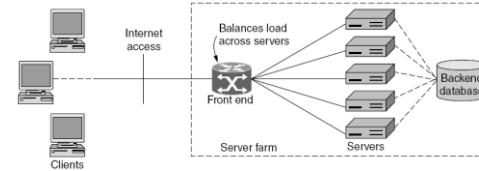
To make the servers appear to be one Web site

- DNS solution
 - When a DNS request is made for the web URL, the DNS server returns a **rotating list of the IP addresses** of the servers;
 - Each client tries one IP address, typically the first on the list.
- Front end solution
 - To broadcast all of the incoming requests, each server answers them by agreement
 - The front end is maybe a switch or a router (**maybe a NAT**).
 - The front end maps them to a server
 - The front end need inspect the IP, TCP, or HTTP headers of packers, and notice the servers' state (**maybe a host**)
 - Load balance: round-robin



Server Farm

- Issue: No matter how much bandwidth **one machine** has, it can only serve so many Web requests before the load is too great.
Solution: Server farm



- 1) Each server must have a copy of the Web site;
- 2) The servers are connected to a common back-end database;
- 3) How to make up the server farm **look like** a single logical Web site to client?



Content Delivery Networks

- Basic idea: Install a bunch of servers across the internet and simply **replicate** web pages on those servers. Be sure to **redirect** clients to the **nearest** replica server.

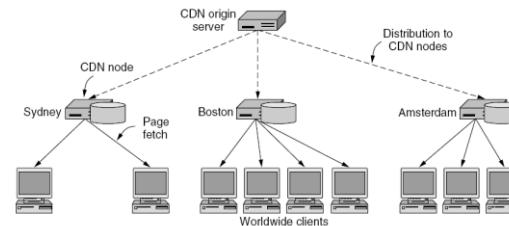
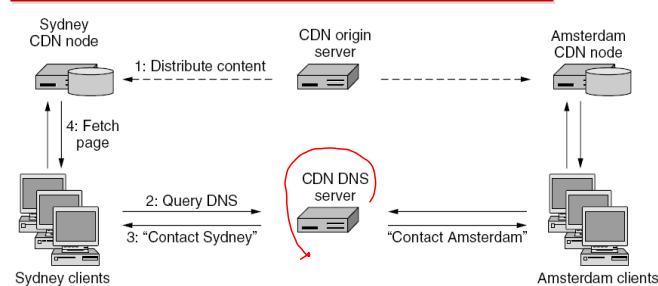


Figure 7-67. CDN distribution tree.

- CDN VS WebProxy = Push vs Pull, public vs private,...



Akamai: Using DNS to redirect...



- CDN DNS server looks at the IP address of the client making the request and returns the IP address of the CDN node that is **nearest** the client. ($\text{nearest} = f(\text{distance}, \text{load})$)



Summary

Our study of network apps now complete!

- Application architectures
 - client-server
 - P2P
 - hybrid
- specific protocols:
 - HTTP
 - DNS
- application service requirements:
 - reliability, bandwidth, delay
- Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP



Summary

Most importantly: learned about protocols

- typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- control vs. data msgs
 - in-band, out-of-band
- centralized vs. decentralized
- stateless vs. stateful
- reliable vs. unreliable msg transfer
- "complexity at network edge"
- message formats:
 - headers: fields giving info about data
 - data: info being communicated

